1. **1 / 1 punto**

The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- ○ The state has multiple numbers rather than only a single number (such as position in the x-direction)
- ○ The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ◉ The state contains numbers such as position and velocity that are continuous valued.
- ○ The reward contains numbers that are continuous valued

✓ **Correcto**
That's right!

2. **1 / 1 punto**

In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is y = _____?

- ○ $y = R(s')$ where s' is the state you get to after taking action a in state s
- ◉ $y = R(s) + \gamma \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ○ $y = \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ○ $y = R(s)$

✓ **Correcto**

3. **1 / 1 punto**

You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

- [x] What an accomplishment -- you made it!

  Correcto

- [x] Andrew sends his heartfelt congratulations to you!

  Correcto

- [x] The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

  Correcto

- [x] You deserve to celebrate!

  Correcto